Introduction to Linear Models

Bert van der Veen

Department of Mathematical Sciences, NTNU

The orchids example

```
set.seed(12345) # For reproducibility
n.times <- 50;n.picks=10;p.orchid <- 0.4
y <- rbinom(n.times, size = n.picks, prob = p.orchid)
```

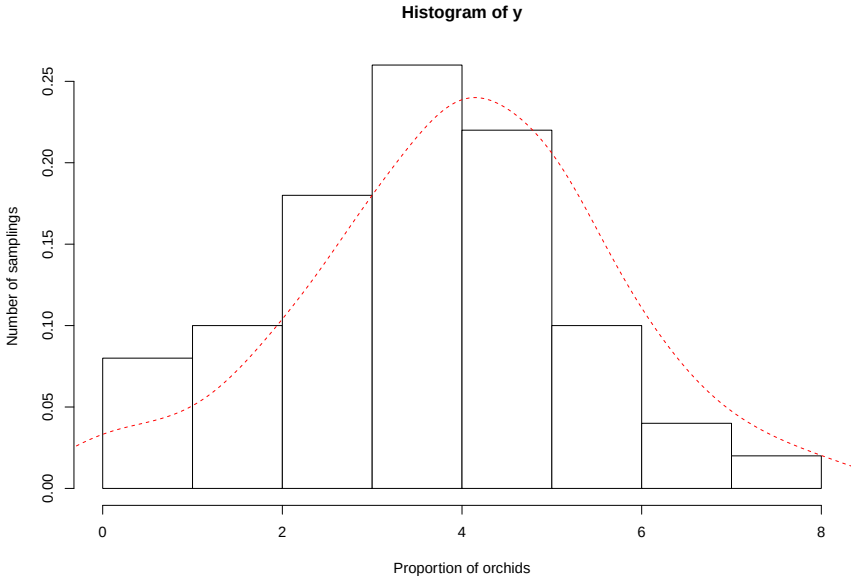

The orchids example

```
set.seed(12345) # For reproducibility
n.times <- 100;n.picks=1;p.orchid <- 0.4
y <- rbinom(n.times, size = n.picks, prob = p.orchid)
```

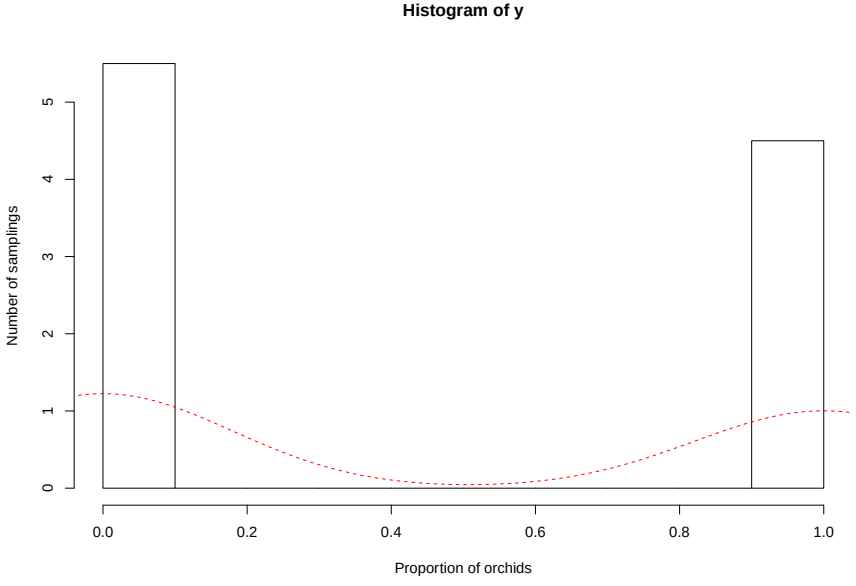

The orchids example

The Binomial distribution is approximately normal for large n_{picks} and π away from 0 and 1

Normality

Real ecological observations are rarely normally distributed.
But it is a nice starting point when learning GLMs.
We can use it (e.g.,) when the mean is far enough from zero.

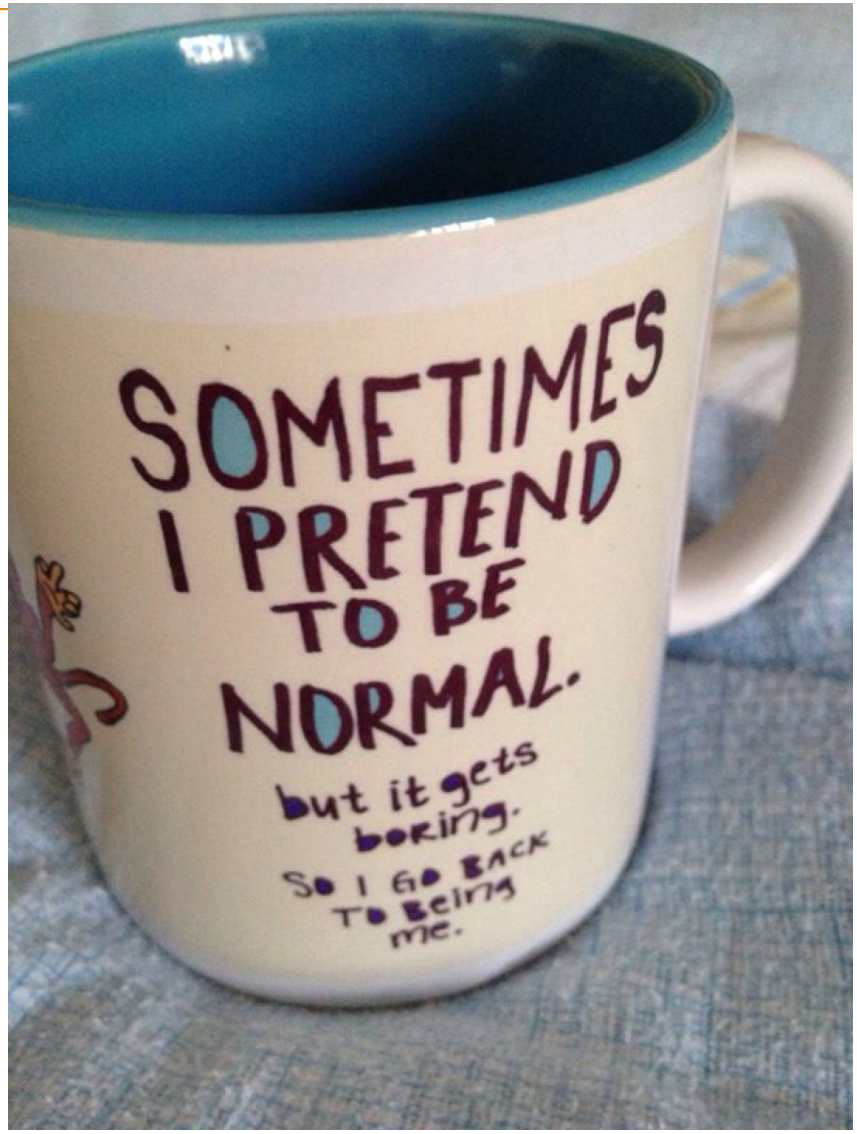

The normal distribution

$$f(y_i; \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left\{ \frac{(y_i - \mu)^2}{2\sigma^2} \right\} \tag{1}$$

The normal distribution (2)

Likelihood:

$$f(y_i; \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left\{ \frac{(y_i - \mu)^2}{2\sigma^2} \right\} \tag{2}$$

log-Likelihood:

$$\log\{f(y_i; \mu, \sigma)\} = -\frac{1}{2} \log(\sigma^2 2\pi) - \frac{(y_i - \mu)^2}{2\sigma^2} \tag{3}$$

Two parameters: μ and σ

- μ is the mean; the middle of the distribution
- σ is the standard deviation; it controls the width

Not only used for data, also the basis of many statistics (e.g., asymptotic sampling distributions)

Estimating μ

Same process as before: calculate gradient and find estimator

$$\frac{\partial \log\{\mathcal{L}(\mathbf{y}; \hat{\mu}, \sigma)\}}{\partial \mu} = \frac{1}{2\sigma^2}\left(2\sum_{i=1}^{n} y_i - 2n\mu\right) \tag{4}$$

Giving..

$$\hat{\mu} = \frac{1}{n}\sum_{i=1}^{n} y_i \tag{5}$$

It is a linear function of y_i so is normally distributed.

Estimating σ^2

Same process as before. MLE is biased so gets a small correction.

$$\hat{\sigma}^2 = \frac{1}{n-1} \sum_{i=1}^{n} (y_i - \hat{\mu})^2 \tag{6}$$

Is it a quadratic function of y_i so is χ^2-distributed.

Uncertainty of $\hat{\mu}$

$$\begin{aligned} \text{var}(\hat{\mu}) &= \mathbb{E}(\hat{\mu}^2) - \mathbb{E}(\hat{\mu})\mathbb{E}(\hat{\mu}) \\ &= \frac{1}{n^2}\sum_{i=1}^{n}\mathbb{E}(y_i^2) - \mathbb{E}(\hat{\mu})\mathbb{E}(\hat{\mu}) \\ &= \frac{1}{n}\sigma^2 \end{aligned} \tag{7}$$

- Depends on n (small n, large uncertainty)
- Depends on σ^2
- Which is estimated by $\hat{\sigma}^2$
- But that estimate also has uncertainty

So, we use the t-distribution to represent that additional uncertainty.

The t-distribution

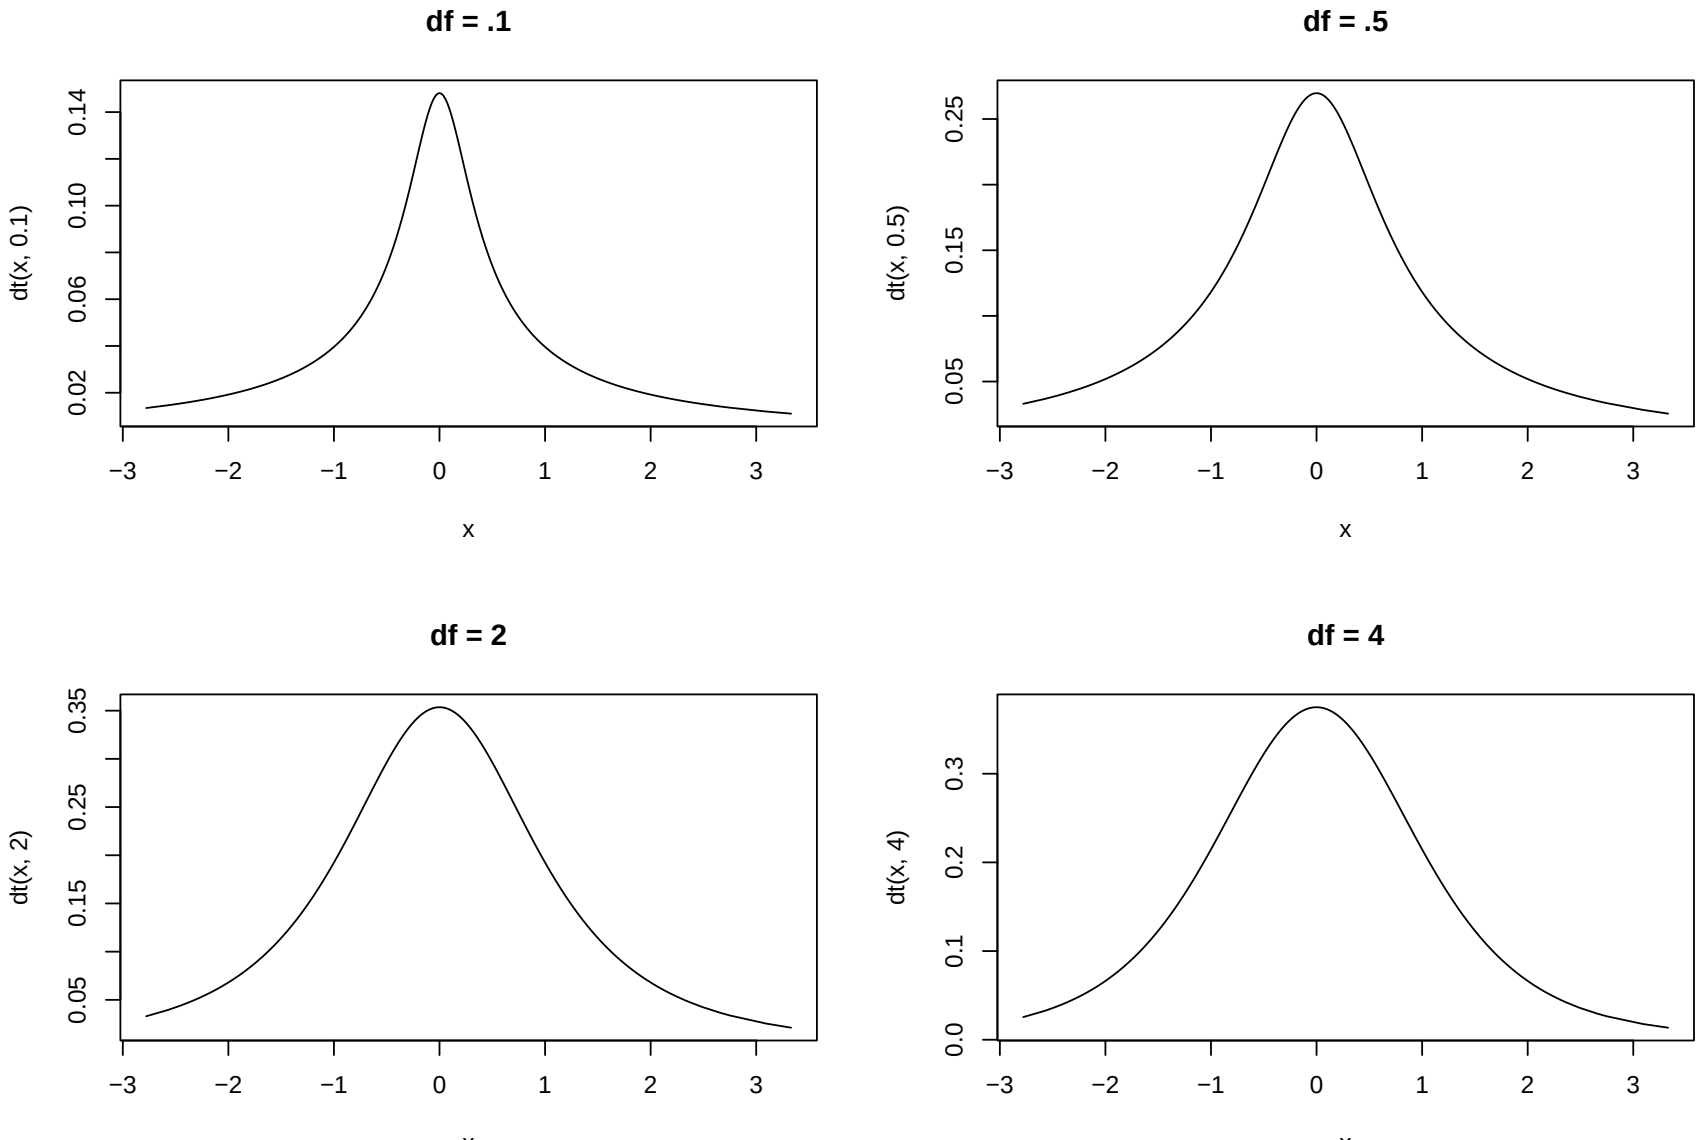

The t-test

- The t-test is type of linear regression
- With one covariate and two groups

```
set.seed(12345)
y <- rnorm(10)
x <- rnorm(10, mean = 2)
t.test(x, y)
```

```
##
##  Welch Two Sample t-test
##
## data:  x and y
## t = 6.5336, df = 17.979, p-value = 3.872e-06
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  1.641042 3.196802
## sample estimates:
##  mean of x  mean of y
##  2.2859778 -0.1329441
```

t-test visualized

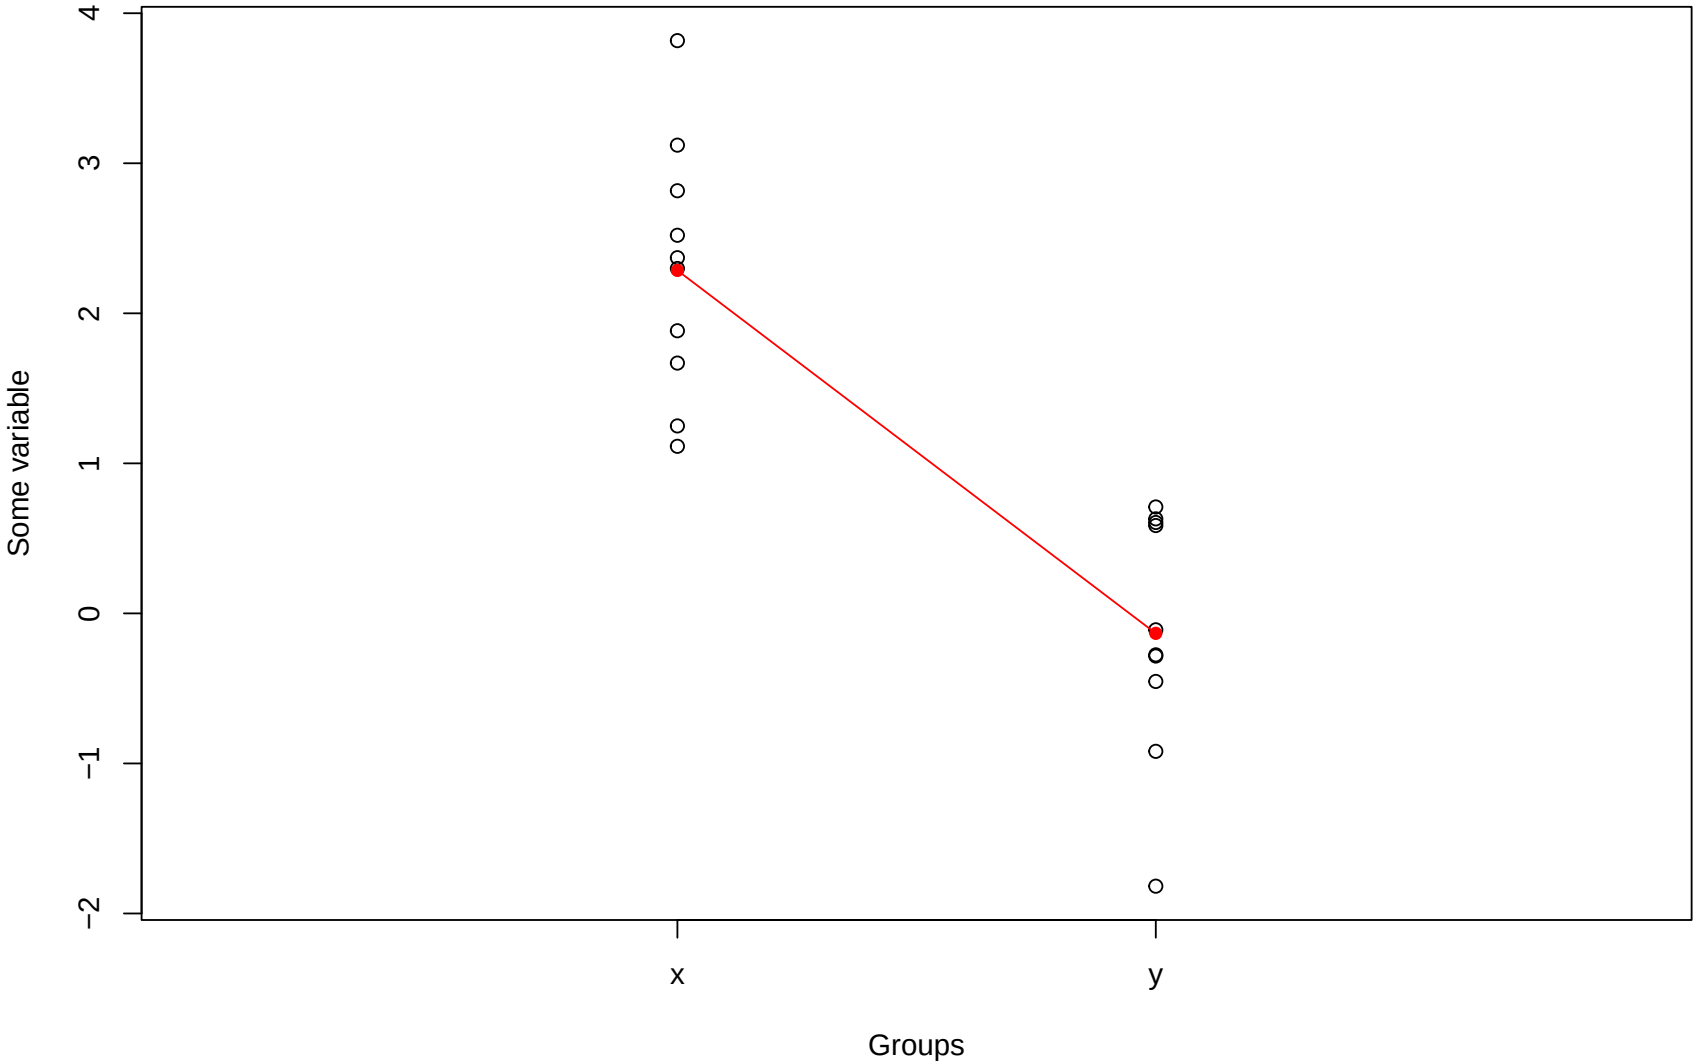

t-test as linear regression

```
data <- data.frame(y = c(x, y),
                   var = c(rep(c("x","y"),each=10)))
lm(y~0+var, data = data)
```

```
##
## Call:
## lm(formula = y ~ 0 + var, data = data)
##
## Coefficients:
##    varx     vary
##  2.2860  -0.1329
```

What is a linear regression?

Models with a continuous **response variable** as a function of one or more **explanatory variable**. Variables are connected by linear equations.

- y_i: the **response variable**, can only be numerical
- x_i: the **explanatory variable**, can be categorical (0,1) or numerical

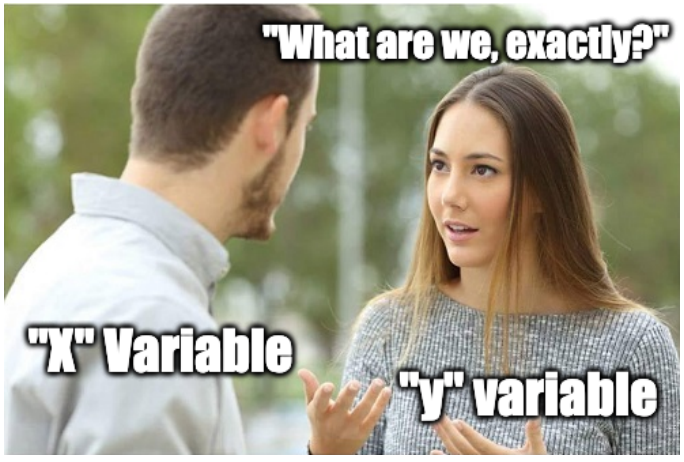

$$y_i = \alpha + x_i\beta + \epsilon_i \sim \mathcal{N}(0, \sigma^2) \tag{8}$$

Synonyms

- Covariate
- Predictor (variable)
- Explanatory variable
- Independent variable

They all refer to x_i.

What is the goal of regression?

We measure data y_i and want to infer its with x_i

Steps:

1) We decide on a model
2) We estimate the parameters
3) We check if it is a valid and good model
4) We draw our conclusion (with uncertainty)

Examples of linear models: categorical x_i

$$\mu_i = \begin{cases} \beta_0 & \text{if } X_i = 0 \\ \beta_1 & \text{if } X_i = 1 \end{cases} \qquad y_i = (1 - x_i)\beta_0 + x_i\beta_1 + \epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

$$\mu_i = \begin{cases} \alpha & \text{if } X_i = 0 \\ \alpha + \beta & \text{if } X_i = 1 \end{cases} \qquad y_i = \alpha + x_i\beta_1 + \epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

Examples of linear models: categorical x_i

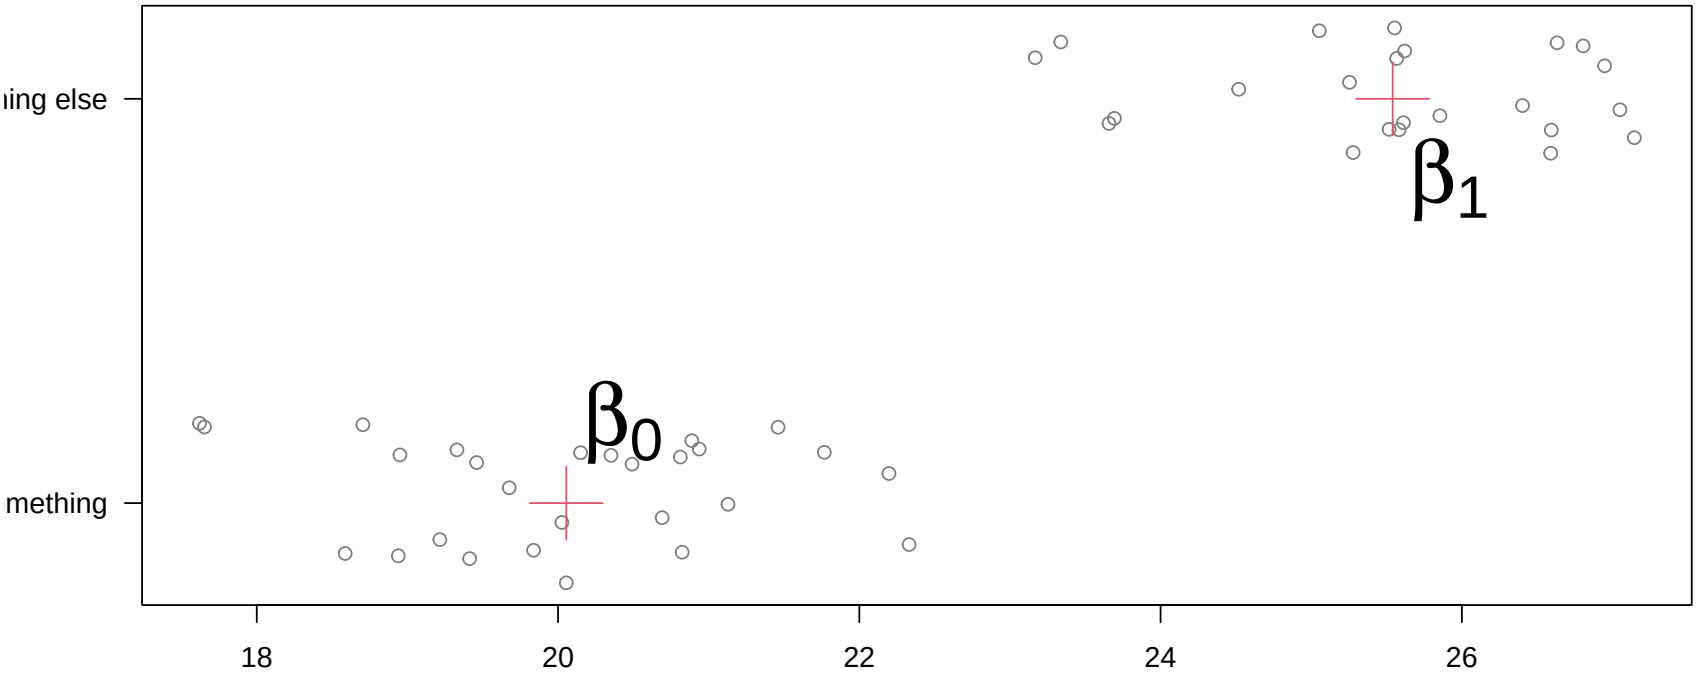

- β_0 is the group 1 mean
- β_1 is the group 2 mean

Examples of linear models: categorical x_i

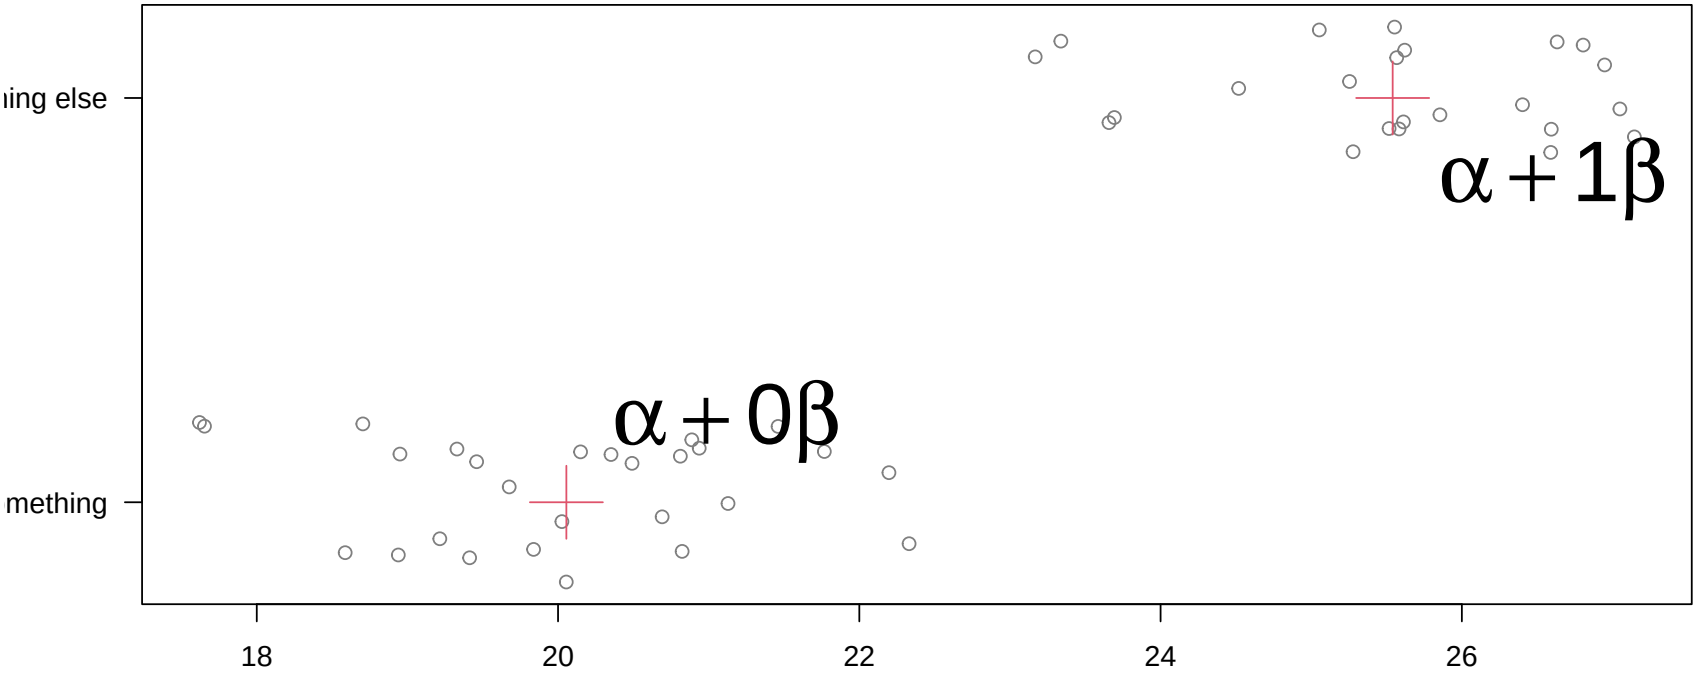

- α is the mean of the first group
- β is the deviation from the mean of the first group

Examples of linear models: continuous x_i

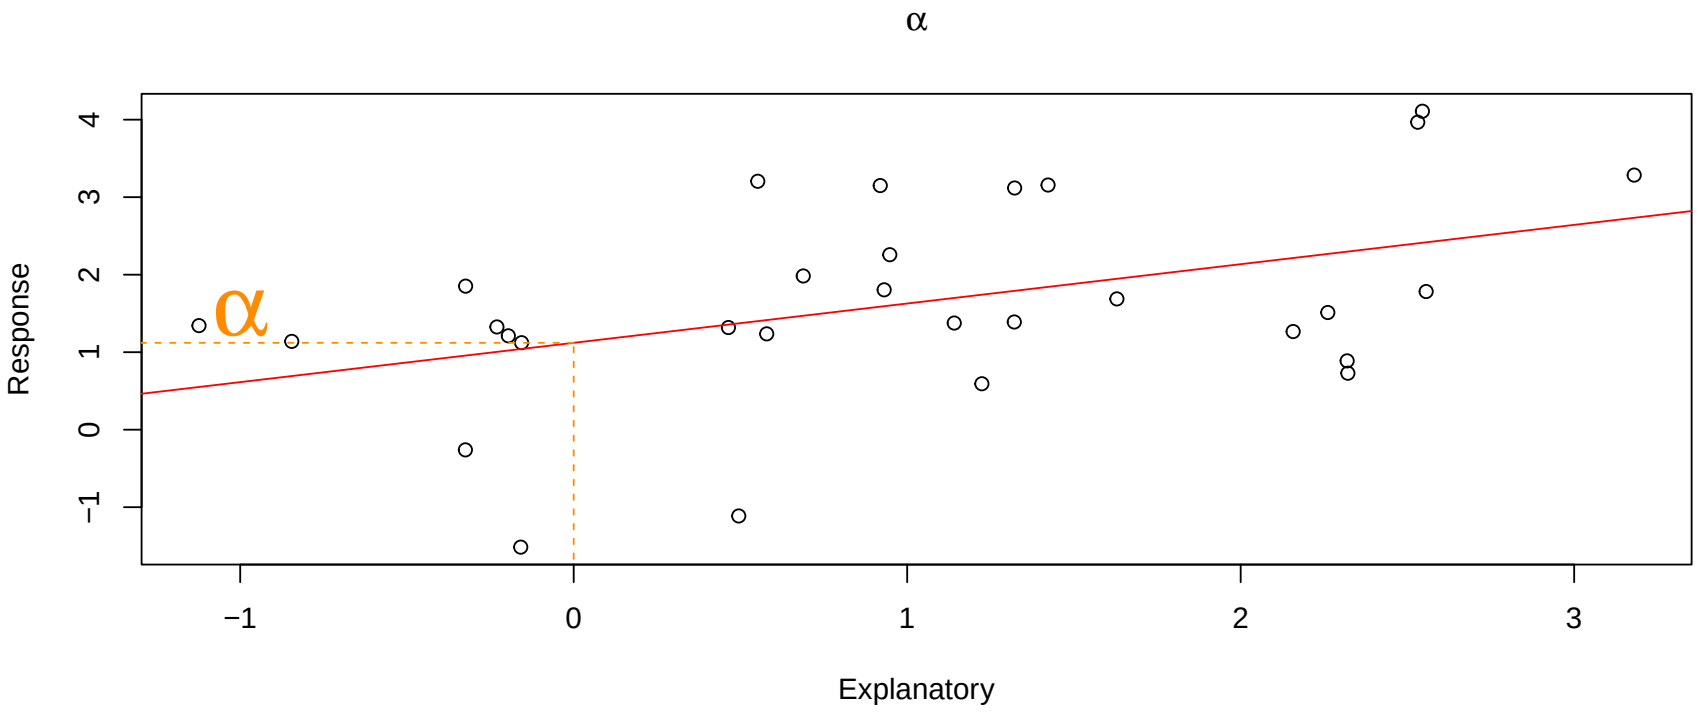

$$y_i = \alpha + x_i\beta + \epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

- α: the intercept is the value of y_i where $x_i = 0$

Examples of linear models: continuous x_i

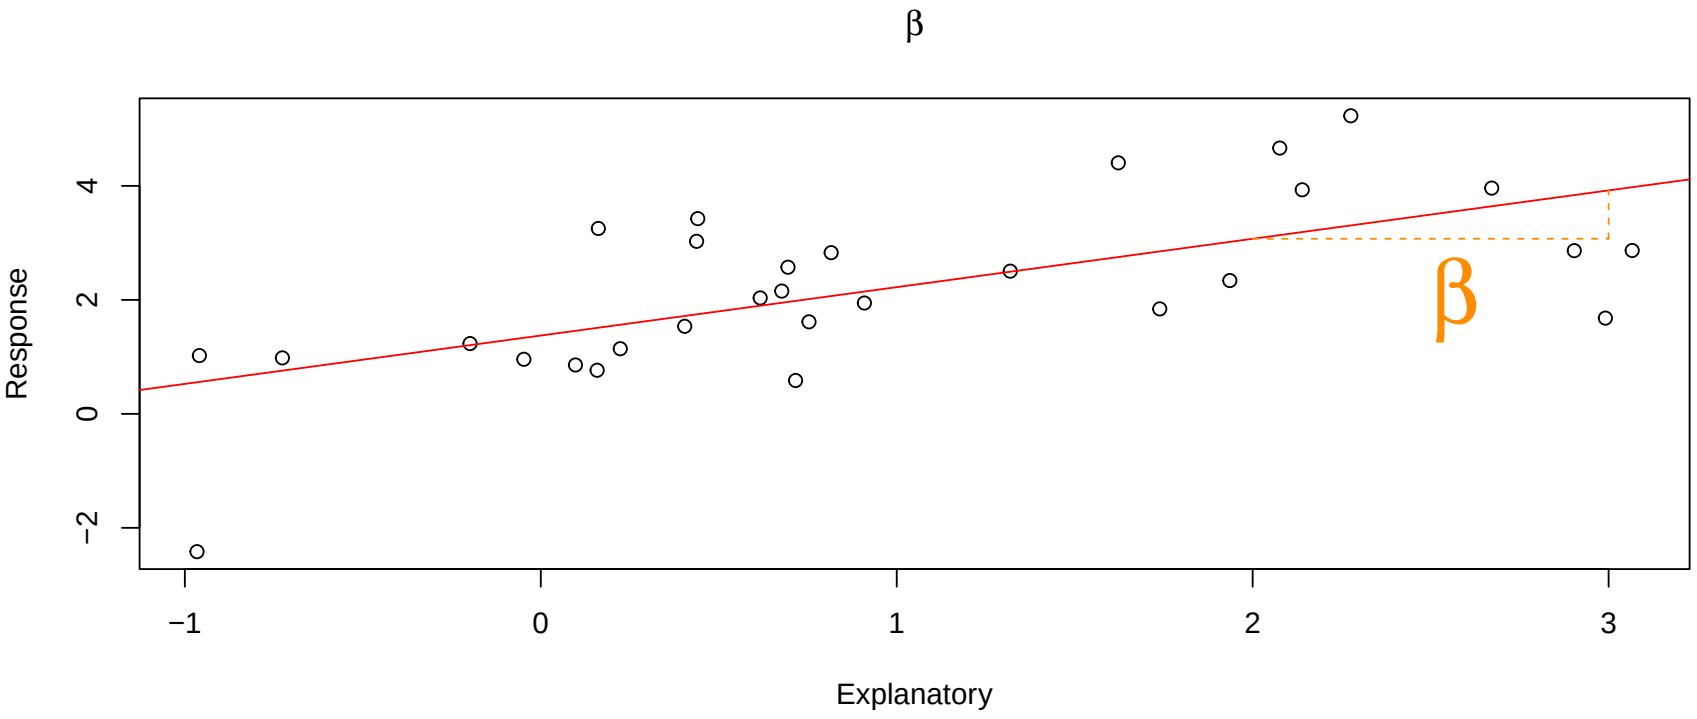

$$y_i = \alpha + x_i\beta + \epsilon_i \sim \mathcal{N}(0, \sigma^2)$$

- α: the intercept is the value of y_i where $x_i = 0$
- β: the slope is the change in y_i for a unit increase in x_i

What is the best line?

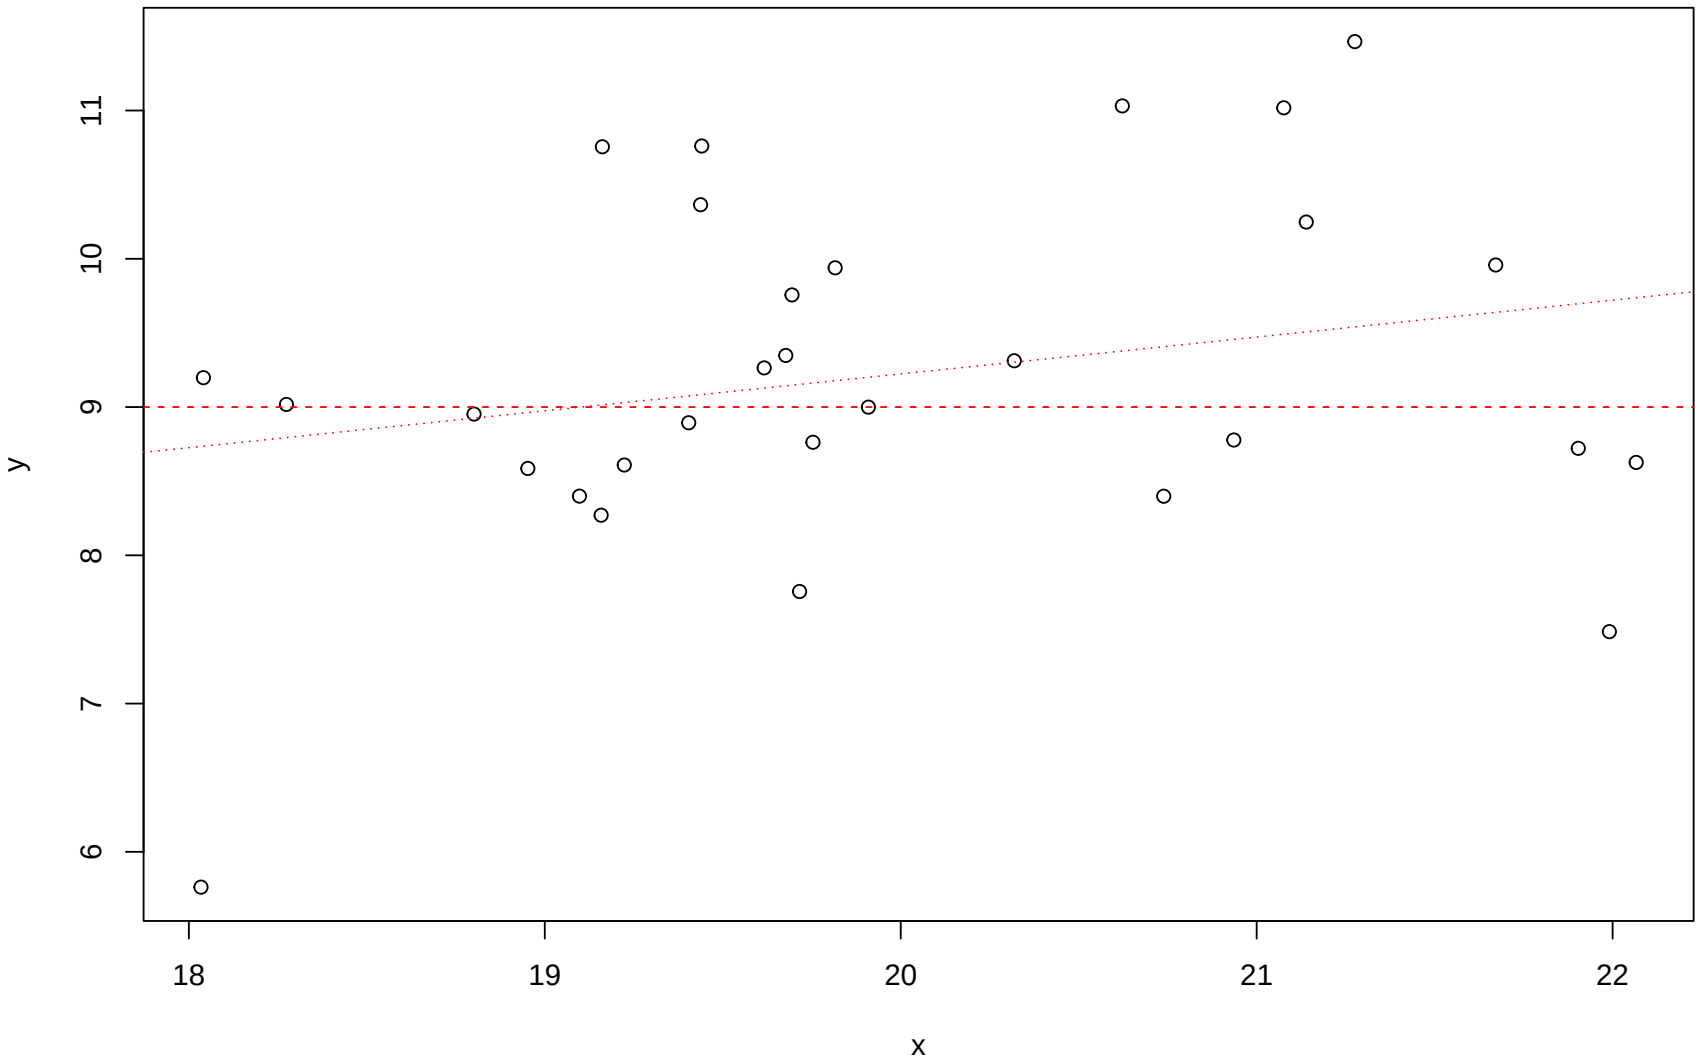

How good is the line?

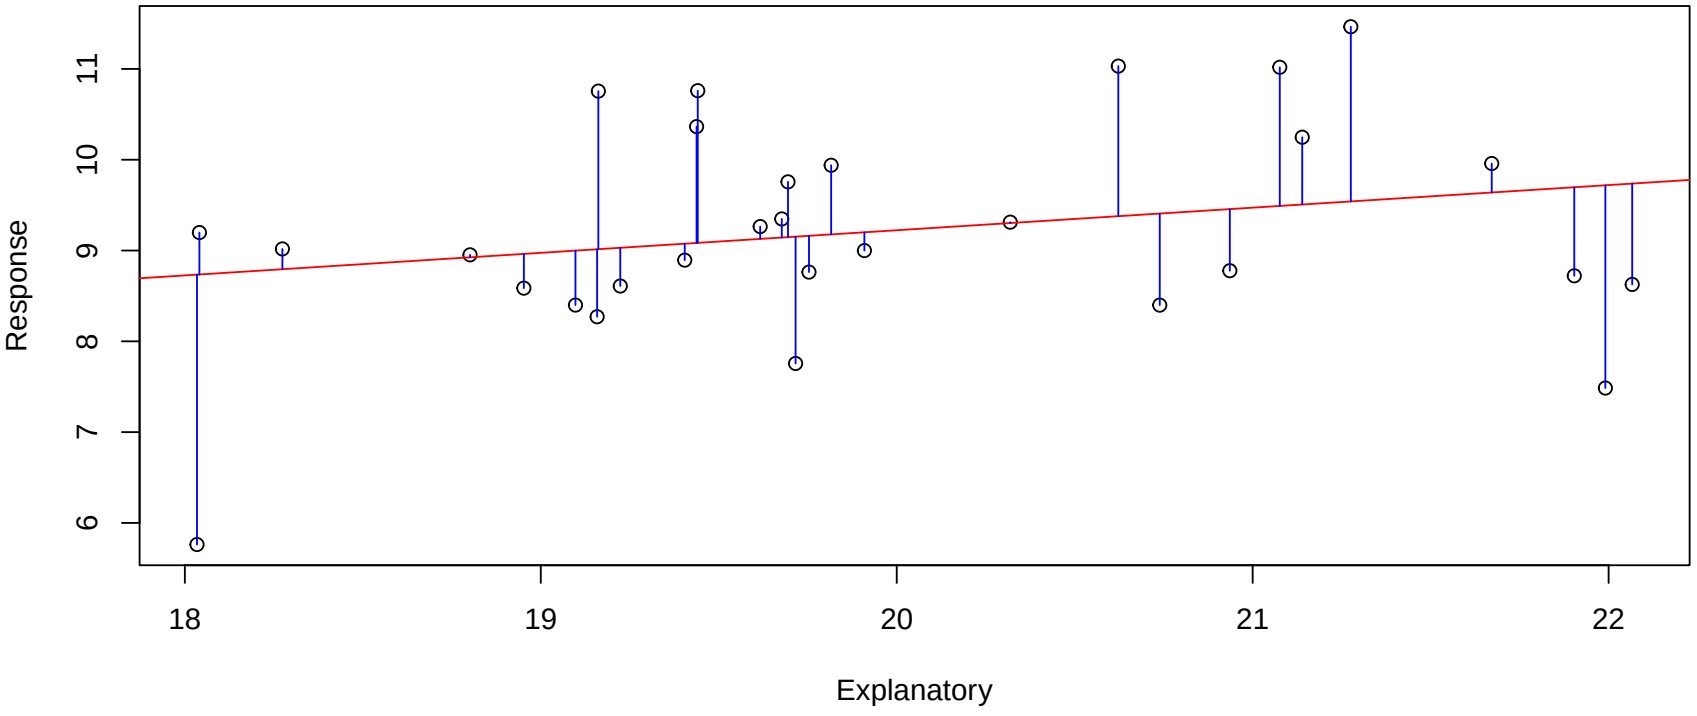

Distance from model (line) to data: "error" ϵ_i

Least squares estimation

Minimize the sum of squared residuals:

$$RSS = \sum_{i=1}^{n} \epsilon_i^2 \tag{9}$$

which is the same to maximizing the normal likelihood!

- $\hat{\alpha} = \frac{1}{N}(\sum y_i - \hat{\beta} \sum x_i)$
- $\hat{\beta} = \frac{\sum (x_i - \sum \frac{x_i}{N})(y_i - \sum \frac{y_i}{N})}{\sum (x_i - \sum \frac{x_i}{N})^2}$
- $\hat{\sigma^2} = \frac{1}{N-1} \sum (y_i - (\alpha + x_i \beta))^2$
- $\hat{\mu}_i = \hat{\alpha} + x_i \hat{\beta}$

Our model

$$y_i = \alpha + \beta x_i + \epsilon_i \sim \mathcal{N}(0, \sigma^2) \tag{10}$$

- y_i: our data
- α, β describe our line
- ϵ_i quantifies distance to the model

Assumptions

We make some critical assumptions here

1) the relationship between y_i and x_i
2) distribution of the errors

Other assumptions for the Errors

$$\epsilon_i \overset{iid}{\sim} \mathcal{N}(0, \sigma^2) \tag{11}$$

- are normally distributed
- have constant variance ("Homoscedastcity")
- are independent
- no outliers

Assumptions for the errors

all the errors together tell us how good the line is

- this is the same as finding the line by maximum likelihood estimation 😄

$$y_i \sim \mathcal{N}(\alpha + x_i\beta, \sigma^2) \tag{12}$$

Summary simple linear models

- includes t-test, anova (analysis of variance) and regression
- all use the same mathy bits (and model)
- **interpretation** depends on the type of **variable**
- GLMs take the same form

Summary

- Normal-distribution and t-distribution
- Simple linear regression: one covariate
- Least squares estimation
- Difference in LMs with categorical or continuous explanatory variables
  - Categorical: intercept/mean parameter
  - Continuous: slope parameter
- Fortunately we have the `lm()` function in R!
- More on assumptions checking tomorrow

Summary (2)

t-test: linear model with categorical covariate of two groups
ANOVA: linear model with categorical covariate of multiple groups
ANCOVA: linear model with categorical and continuous covariate (interaction)